



**Faculty of Information Technology**  
**University of Moratuwa**  
**BSc Hons in Information Technology**  
**BSc Hons in Information Technology Management**  
**IN 4410 – Big Data Analytics**

**Level 1 - Semester 2**

**Lab Sheet 08**

## **Spark SQL**

(for Windows)

Spark introduces a programming module for structured data processing called Spark SQL. It provides a programming abstraction called Data Frame and can act as a distributed SQL query engine.

Unlike the basic Spark RDD API, the interfaces provided by Spark SQL provide Spark with more information about the structure of both the data and the computation being performed. Internally, Spark SQL uses this extra information to perform extra optimizations.

There are several ways to interact with Spark SQL including SQL and the Dataset API. When computing a result, the same execution engine is used, independent of which API/language you are using to express the computation. This unification means that developers can easily switch back and forth between different APIs based on which provides the most natural way to express a given transformation.

Let's install PySpark on your PC.

### **Pre-Requisites:**

**Method 1** - Jupyter Notebook

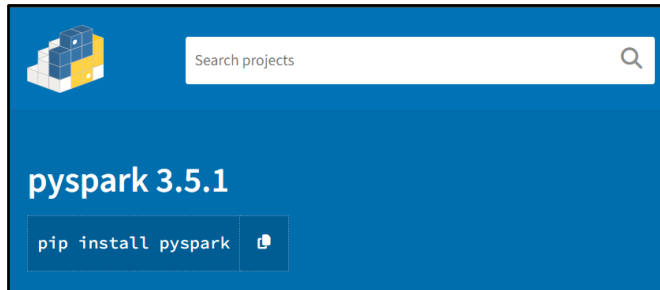
Or

**Method 2** - Databricks Community Edition

You can follow any of these two methods.

## ❖ Jupyter Notebook

- 1) Type Pyspark pypi in the browser and Go to the pyspark web page and copy the command **pip install pyspark**.



- 2) Open a command prompt as administrator and paste the command and install pyspark to the machine.

```
C:\WINDOWS\system32>pip install pyspark
Collecting pyspark
  Downloading pyspark-3.5.1.tar.gz (317.0 MB)
    ----- 316.9/317.0 MB 12.1 MB/s eta 0:00:01
```

- 3) After installing pyspark type python and then import pyspark feature. If there are not any errors, it's successfully installed.

```
C:\WINDOWS\system32>python
Python 3.12.2 (tags/v3.12.2:6abddd9, Feb 6 2024, 21:26:36) [MSC v.1937 64 bit (AMD64)] on win32
Type "help", "copyright", "credits" or "license" for more information.
>>> import pyspark
>>>
```

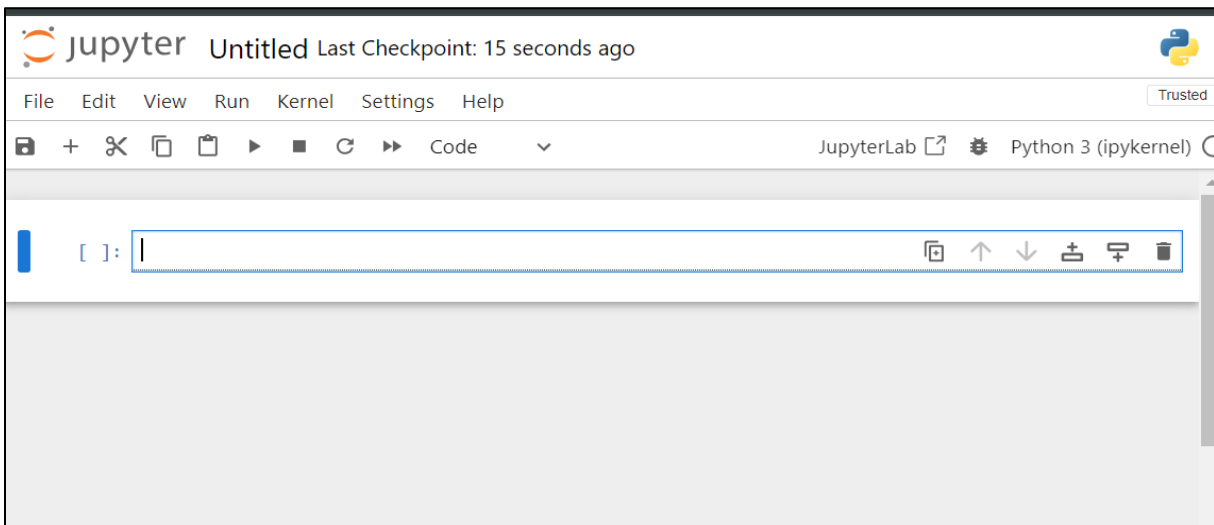
- 4) Then go to the Scripts folder of python and type cmd command in the path, it will open a new command prompt and type pip install jupyter. it will install the jupyter notebook.

```
Downloading types_python_dateutil-2.9.0.20240316-py3-none-any.whl.metadata (1.8 kB)
Downloading jupyter-1.0.0-py2.py3-none-any.whl (2.7 kB)
Downloading ipykernel-6.29.3-py3-none-any.whl (117 kB)
    ----- 117.1/117.1 kB 488.0 kB/s eta 0:00:00
Downloading ipywidgets-8.1.2-py3-none-any.whl (139 kB)
    ----- 139.4/139.4 kB 1.0 MB/s eta 0:00:00
Downloading jupyter_console-6.6.3-py3-none-any.whl (24 kB)
Downloading nbconvert-7.16.3-py3-none-any.whl (257 kB)
```

- 5) Create a new folder in any drive name called python and open that with cmd and type jupyter notebook. It will open a new jupyter notebook in browser.

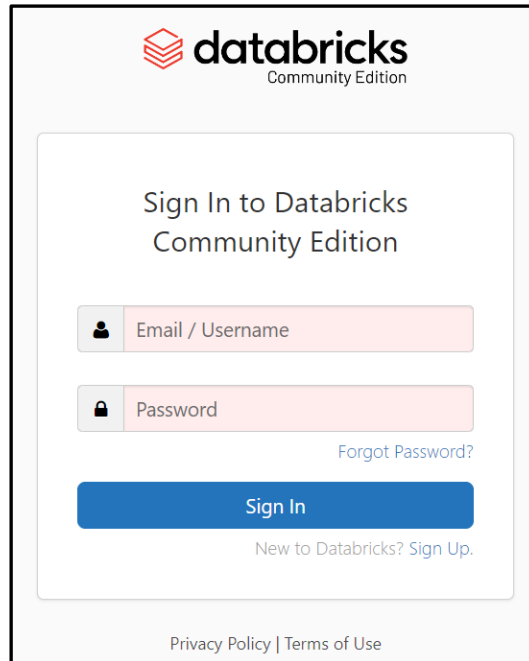
```
C:\Users\MMRL-Acer\Documents\Python>jupyter notebook
[I 2024-03-26 12:01:47.002 ServerApp] Extension package jupyter_lsp took 0.1055s to import
[I 2024-03-26 12:01:47.487 ServerApp] Extension package jupyter_server_terminals took 0.4726s to import
[I 2024-03-26 12:01:47.846 ServerApp] jupyter_lsp | extension was successfully linked.
[I 2024-03-26 12:01:47.846 ServerApp] jupyter_server_terminals | extension was successfully linked.
[I 2024-03-26 12:01:47.862 ServerApp] jupyterlab | extension was successfully linked.
[I 2024-03-26 12:01:47.877 ServerApp] notebook | extension was successfully linked.
[I 2024-03-26 12:01:47.877 ServerApp] Writing Jupyter server cookie secret to C:\Users\MMRL-Acer\AppData\Roaming\jupyter
\runtime\jupyter_cookie_secret
```

- 6) Then you can open the notebook and rename the shell as any name and select the python as the type and then you can use it.

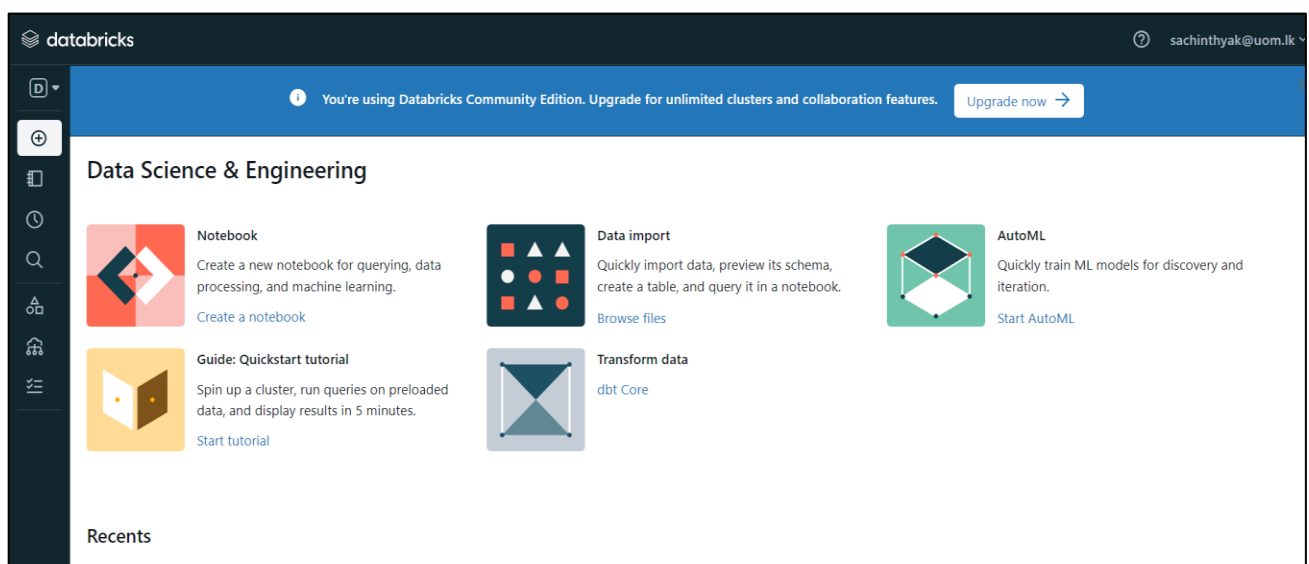


## ❖ Databricks Community Edition

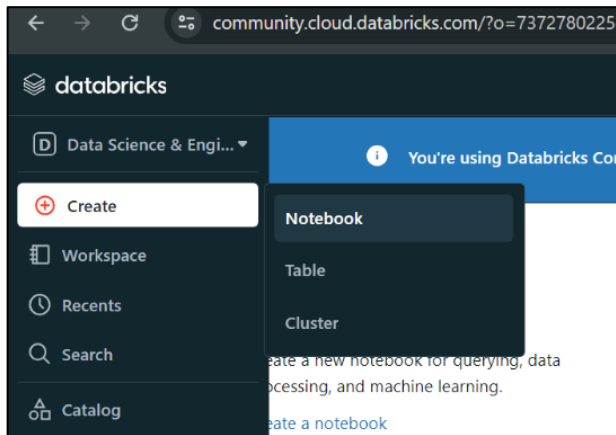
- 1) Create a new account in Databricks and sign in.



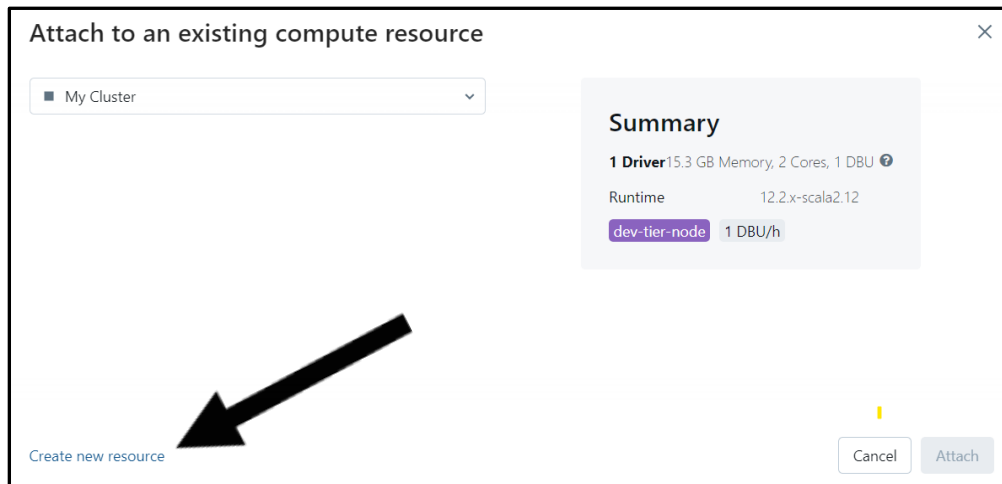
- 2) Open the Databricks cloud and rename the workspace as any name and select python as the type.



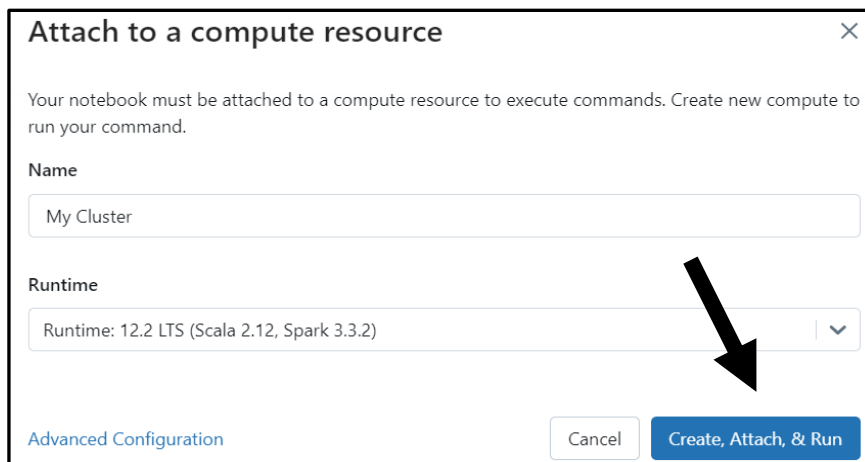
3) Create a new **Notebook** to start coding.



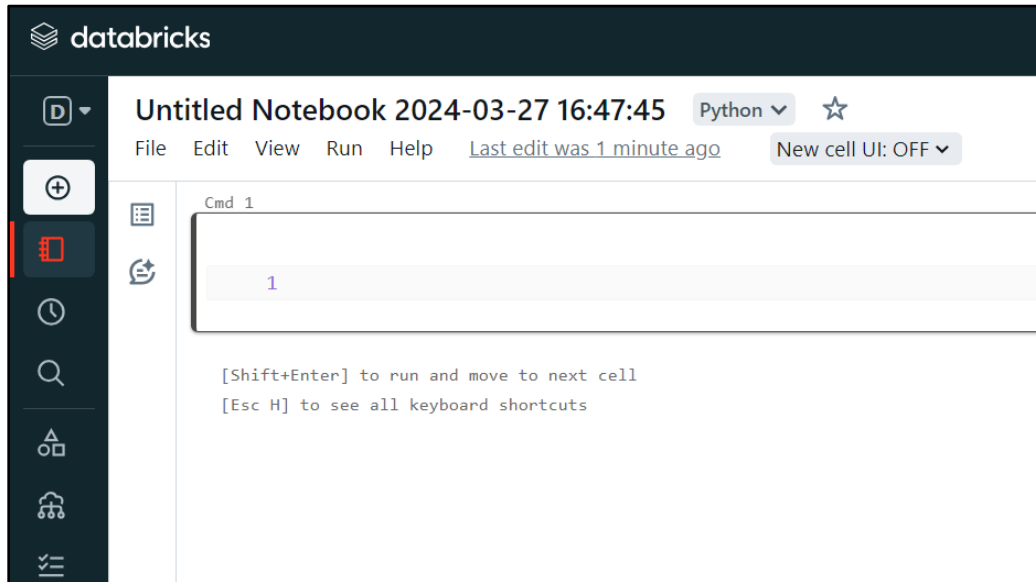
4) Create a new cluster for the analysis, press **Create new resource** tab.



5) Press **Create, Attach & Run** tab to create the cluster.



- 6) Finally, it will appear a window as below as same as jupyter notebook then you can use it.



## 1. Create a PySpark RDD

- Import and create a SparkSession:

```
from pyspark.sql import SparkSession
spark = SparkSession.builder.getOrCreate()
```

- Create a rdd:

```
rdd = spark.sparkContext.parallelize([("Colombo", 1),
("Kalutara", 2), ("Gampaha", 3)])
```

```
rdd
```

```
1 rdd
```

```
Out[6]: ParallelCollectionRDD[1] at readRDDFromInputStream at PythonRDD.scala:435
```

```
rdd.collect()
```

```
1 rdd.collect()
```

► (1) Spark Jobs

```
Out[7]: [('Colombo', 1), ('Kalutara', 2), ('Gampaha', 3)]
```

```
rdd.count()
```

```
1 rdd.count()
```

► (1) Spark Jobs

```
Out[9]: 3
```

## 2. Create a PySpark Data Frame from RDD

- Import and create a Spark Session:

```
from pyspark.sql import SparkSession
from datetime import date, datetime
spark = SparkSession.builder.getOrCreate()
```

```
rdd = spark.SparkContext.parallelize([
    (1, 1.0, "string1", date(2024, 1, 1), datetime(2024, 1, 1, 12, 0)),
    (2, 2.0, "string2", date(2024, 2, 1), datetime(2024, 2, 1, 12, 0)),
    (3, 3.0, "string3", date(2024, 3, 1), datetime(2024, 3, 1, 12, 0))
])

df=spark.createDataFrame(rdd, schema=["num", "float", "string", "date", "datetime"])
```

df

```
1 df = spark.createDataFrame(rdd, schema= ["num","float","string","date","datetime"])
```

▼ (2) Spark Jobs

▶ Job 2 [View](#) (Stages: 1/1)

▶ Job 3 [View](#) (Stages: 1/1)

▼ df: pyspark.sql.dataframe.DataFrame

num: long

float: double

string: string

date: date

datetime: timestamp

df.show()

```
1 df.show()
```

▶ (3) Spark Jobs

num	float	string	date	datetime
1	1.0	string1	2024-01-01	2024-01-01 12:00:00
2	2.0	string2	2024-02-01	2024-02-01 12:00:00
3	3.0	string3	2024-03-01	2024-03-01 12:00:00

df.show(1)

```
1 df.show(1)
```

▶ (3) Spark Jobs

num	float	string	date	datetime
1	1.0	string1	2024-01-01	2024-01-01 12:00:00

only showing top 1 row



```
df.printSchema()
```

```
1 df.printSchema()  
  
root  
 |-- num: long (nullable = true)  
 |-- float: double (nullable = true)  
 |-- string: string (nullable = true)  
 |-- date: date (nullable = true)  
 |-- datetime: timestamp (nullable = true)
```

### 3. Create a Spark context in Python

- Import and create a Spark Session, Spark Context and SQL Context:

```
import pyspark  
from pyspark import SparkContext  
from pyspark.sql import SparkSession  
from pyspark.sql import SQLContext
```

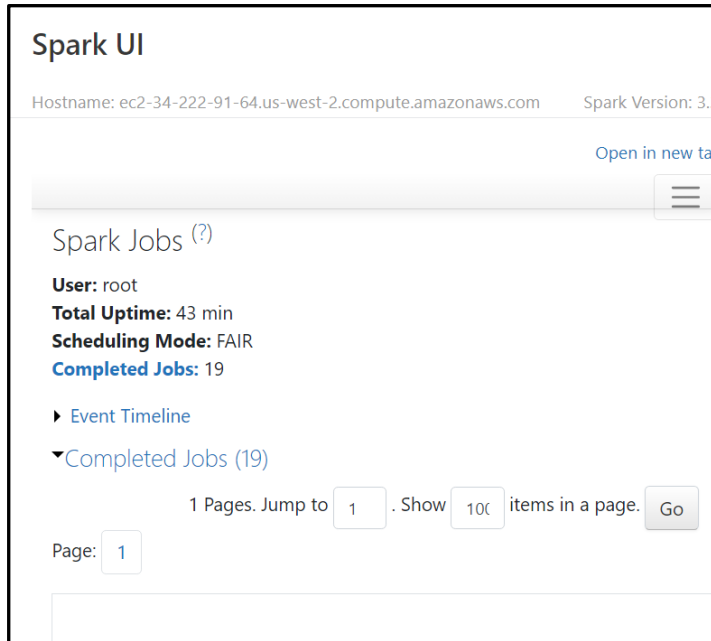
- Memory allocation for spark Context

```
conf = pyspark.SparkConf().setMaster("local").setAppName("My  
First Spark Practical").setAll([("spark.driver.memory", "40g"),  
("spark.executor.memory", "50g")])
```

```
sc = SparkContext.getOrCreate(conf=conf)
```

```
1 sc = SparkContext.getOrCreate(conf=conf)  
  
Command took 0.13 seconds -- by sachinthyak@uom.lk at 3/28  
md 4  
  
1 sc  
  
SparkContext  
Spark UI  
Version  
v3.3.2  
Master  
local[8]  
AppName  
Databricks Shell
```

After creating the spark context, simply click the [Spark UI](#) link and then you can move on to the Spark UI Shell



To stop the spark context, you can type,

```
sc.stop()
```

## 4. Create a PySpark Data Frame

- Import and create a SparkSession:

```
from pyspark.sql import SparkSession
spark = SparkSession.builder.getOrCreate()
from datetime import date
```

- Create data frame

```
df = spark.createDataFrame([
    ("Red", 1, "Apple", date(2024,1,1)),
    ("Black", 2, "Grapes", date(2024,2,1)),
    ("Yellow", 3, "Banana", date(2024,3,1))],
    schema = "color string, sr_no long, fruit string, datet date")
```

```
df: pyspark.sql.dataframe.DataFrame
  color: string
  sr_no: long
  fruit: string
  datet: date
```

df

```
1 df
```

```
Out[5]: DataFrame[color: string, sr_no: bigint, fruit: string, datet: date]
```

df.show

```
+-----+-----+-----+-----+
| color|sr_no| fruit|    datet|
+-----+-----+-----+-----+
|  Red|    1| Apple|2024-01-01|
| Black|    2| Grapes|2024-02-01|
|Yellow|    3| Banana|2024-03-01|
+-----+-----+-----+-----+
```


df.show(2)

► (3) Spark Jobs

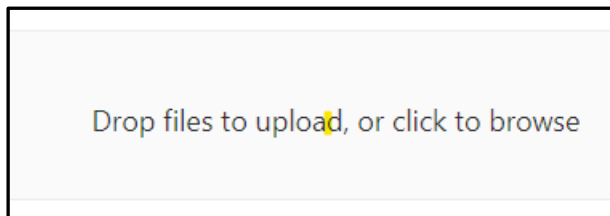
```
+-----+-----+-----+-----+
| color|sr_no| fruit|    datet|
+-----+-----+-----+-----+
|  Red|    1| Apple|2024-01-01|
|Black|    2| Grapes|2024-02-01|
+-----+-----+-----+-----+
only showing top 2 rows
```

## 5. Read a CSV file in Pyspark


Go to the **Catalog** tab in left side corner and click the button



Press the **Create table** button, and then browse or drop the csv file.



You can see it successfully uploaded.



Then press **Create a table in notebook**, it will automatically generate a DBFS system and command prompt area.

```
%fs  
ls /FileStore/tables/iris-6.csv
```

```

1 %fs
2 ls /FileStore/tables/iris-6.csv

```

Table ▾ +				
	path	name	size	modificationTime
1	dbfs:/FileStore/tables/iris-6.csv	iris-6.csv	3975	1711692426000

↓ 1 row | 12.66 seconds runtime

Now, you can load your csv file into your DBFS system

```

1 display(dbutils.fs.ls("FileStore/tables/iris-6.csv"))

```

▼ (3) Spark Jobs

- ▶ Job 0 [View](#) (Stages: 1/1)
- ▶ Job 1 [View](#) (Stages: 1/1)
- ▶ Job 2 [View](#) (Stages: 1/1)

Table ▾ +				
	path	name	size	modificationTime
1	dbfs:/FileStore/tables/iris-6.csv	iris-6.csv	3975	1711692426000

```

irisdatadf =
spark.read.format("csv").option("header","true").option("inferSchema", "true").load("/FileStore/tables/iris-6.csv")

```

▼ (2) Spark Jobs

- ▶ Job 3 [View](#) (Stages: 1/1)
- ▶ Job 4 [View](#) (Stages: 1/1)

▼ irisdatadf: pyspark.sql.dataframe.DataFrame

```

sepal.length: double
sepal.width: double
petal.length: double
petal.width: double
variety: string

```

```
irisdatadf.show()
```

```
1 irisdatadf.show()
```

► (1) Spark Jobs

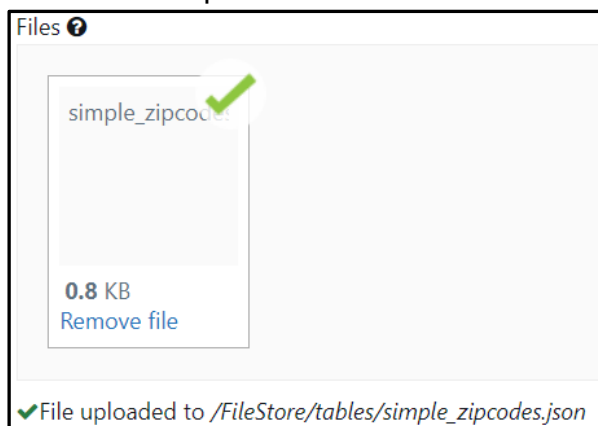
sepal.length	sepal.width	petal.length	petal.width	variety
5.1	3.5	1.4	0.2	Setosa
4.9	3.0	1.4	0.2	Setosa
4.7	3.2	1.3	0.2	Setosa
4.6	3.1	1.5	0.2	Setosa
5.0	3.6	1.4	0.2	Setosa
5.4	3.9	1.7	0.4	Setosa

## 6. Read a JSON file in Pyspark

- Type the code to see the Json methods

```
help(spark.read.json)
```


- Browse or drop the csv file



- Create data frame for json file

```
df =  
spark.read.json(path="/FileStore/tables/simple_zipcodes.json")
```

► (1) Spark Jobs

▼  df: pyspark.sql.dataframe.DataFrame

```
City: string
State: string
ZipCodeType: string
Zipcode: long
```

```
df.printSchema()
```

```
root
|-- City: string (nullable = true)
|-- State: string (nullable = true)
|-- ZipCodeType: string (nullable = true)
|-- Zipcode: long (nullable = true)
```

```
df.show()
```

City State ZipCodeType Zipcode			
PARC PARQUE	PR	STANDARD	704
PASEO COSTA DEL SUR	PR	STANDARD	704
BDA SAN LUIS	PR	STANDARD	709
CINGULAR WIRELESS	TX	UNIQUE	76166
FORT WORTH	TX	STANDARD	76177
FT WORTH	TX	STANDARD	76177
URB EUGENE RICE	PR	STANDARD	704
MESA	AZ	STANDARD	85209
MESA	AZ	STANDARD	85210
HILLIARD	FL	STANDARD	32046

## 7. Create Or Replace Temporary View in Data Frame using PySpark

```
spark.sql("select * from df")
```

If you just type this code only, it will occur an error, so we have to change the code with this `df.createOrReplaceTempView` function so it will create a temporary view to this SQL session. Once you create this, you will be able to execute any SQL query on top of your data frame inside the spark session.

```
df.createOrReplaceTempView("df")
spark.sql("select * from df")
```

```
1 df.createOrReplaceTempView("df")

Command took 0.22 seconds -- by sachinthyak@uom.lk at 3/29/2024, 6:40:53 PM on My Cluster

Cmd 7

1 spark.sql("select * from df")

Out[5]: DataFrame[City: string, State: string, ZipCodeType: string, Zipcode: bigint]
```

```
spark.sql("select * from df").show()
```

```
1 spark.sql("select * from df").show()

▶ (1) Spark Jobs

+-----+-----+-----+-----+
|          City|State|ZipCodeType|Zipcode|
+-----+-----+-----+-----+
|    PARC PARQUE|PR|    STANDARD|    704|
|PASEO COSTA DEL SUR|PR|    STANDARD|    704|
|    BDA SAN LUIS|PR|    STANDARD|    709|
| CINGULAR WIRELESS|TX|    UNIQUE|  76166|
|    FORT WORTH|TX|    STANDARD|  76177|
|    FT WORTH|TX|    STANDARD|  76177|
|    URB EUGENE RICE|PR|    STANDARD|    704|
|          MESA|AZ|    STANDARD|  85209|
|          MESA|AZ|    STANDARD|  85210|
|    HILLIARD|FL|    STANDARD|  32046|
+-----+-----+-----+-----+
```



## 8. Create a Global Temporary View in Data Frame using PySpark

If you just type this select all query only, it will occur an error as same as the first temporary view, so we have to change the code with this `df.createGlobalTempView` function so it will create a temporary view to this SQL session. Once you create this, you will be able to execute any SQL query on top of your data frame inside the spark application.

```
1 spark.sql("select * from global_temp.df")
```

⊞ **AnalysisException: [TABLE\_OR\_VIEW\_NOT\_FOUND]** The table or view 'global\_temp`.`df` cannot be found. Verify the spelling and correctness of the schema and catalog. If you did not qualify the name with a schema, verify the current\_schema() output, or qualify the name with the correct schema and catalog. To tolerate the error on drop use DROP VIEW IF EXISTS or DROP TABLE IF EXISTS.; line 1 pos 14;  
'Project [\*]  
+- 'UnresolvedRelation [global\_temp, df], [], false

```
df.createGlobalTempView("df")  
spark.sql("select * from global_temp.df")
```

```
1 spark.sql("select * from global_temp.df")
```

```
Out[9]: DataFrame[City: string, State: string, ZipCodeType: string, Zipcode: bigint]
```

```
sqlcdf = spark.sql("select * from global_temp.df").show()
```

```
1 sqlcdf = spark.sql("select * from global_temp.df").show()
```

► (1) Spark Jobs

City	State	ZipCodeType	Zipcode
PARC PARQUE	PR	STANDARD	704
PASEO COSTA DEL SUR	PR	STANDARD	704
BDA SAN LUIS	PR	STANDARD	709
CINGULAR WIRELESS	TX	UNIQUE	76166
FORT WORTH	TX	STANDARD	76177
FT WORTH	TX	STANDARD	76177
URB EUGENE RICE	PR	STANDARD	704
MESA	AZ	STANDARD	85209
MESA	AZ	STANDARD	85210
HILLIARD	FL	STANDARD	32046

## 9. PySpark Data Frame Functions

- i) `collect()`
- ii) `take()`
- iii) `count()`
- iv) `select()`
- v) `filter()`
- vi) `sort()`
- vii) `describe()`

```
df.collect()
```

```
Out[11]: [Row(City='PARC PARQUE', State='PR', ZipCodeType='STANDARD', Zipcode=704),
Row(City='PASEO COSTA DEL SUR', State='PR', ZipCodeType='STANDARD', Zipcode=704),
Row(City='BDA SAN LUIS', State='PR', ZipCodeType='STANDARD', Zipcode=709),
Row(City='CINGULAR WIRELESS', State='TX', ZipCodeType='UNIQUE', Zipcode=76166),
Row(City='FORT WORTH', State='TX', ZipCodeType='STANDARD', Zipcode=76177),
Row(City='FT WORTH', State='TX', ZipCodeType='STANDARD', Zipcode=76177),
Row(City='URB EUGENE RICE', State='PR', ZipCodeType='STANDARD', Zipcode=704),
Row(City='MESA', State='AZ', ZipCodeType='STANDARD', Zipcode=85209),
Row(City='MESA', State='AZ', ZipCodeType='STANDARD', Zipcode=85210),
Row(City='HILLIARD', State='FL', ZipCodeType='STANDARD', Zipcode=32046)]
```

Command took 0.68 seconds -- by sachinthyak@uom.lk at 3/29/2024, 7:06:34 PM on My Cluster

```
df.take(2)
```

```
1 df.take(2)
```

► (1) Spark Jobs

```
Out[12]: [Row(City='PARC PARQUE', State='PR', ZipCodeType='STANDARD', Zipcode=704),
Row(City='PASEO COSTA DEL SUR', State='PR', ZipCodeType='STANDARD', Zipcode=704)]
```

```
df.count()
```

```
1 df.count()
```

► (2) Spark Jobs

```
Out[13]: 10
```

```
df.select("City", "State").show()
```

City	State
PARC PARQUE	PR
PASEO COSTA DEL SUR	PR
BDA SAN LUIS	PR
CINGULAR WIRELESS	TX
FORT WORTH	TX
FT WORTH	TX
URB EUGENE RICE	PR
MESA	AZ
MESA	AZ
HILLIARD	FL

```
df.select("City", "State").show(4)
```

City	State
PARC PARQUE	PR
PASEO COSTA DEL SUR	PR
BDA SAN LUIS	PR
CINGULAR WIRELESS	TX

only showing top 4 rows

```
df.show()
```

```
df.filter(df["Zipcode"]>704).show()
```

City	State	ZipCodeType	Zipcode
BDA SAN LUIS	PR	STANDARD	709
CINGULAR WIRELESS	TX	UNIQUE	76166
FORT WORTH	TX	STANDARD	76177
FT WORTH	TX	STANDARD	76177
MESA	AZ	STANDARD	85209
MESA	AZ	STANDARD	85210
HILLIARD	FL	STANDARD	32046

```
df.select("City").filter("City like 'M%' ").show()
```

```
+-----+
|City|
+-----+
|MESA|
|MESA|
+-----+
```

```
df.select("City").filter("City like 'F%' ").show()
```

```
+-----+
|      City|
+-----+
|FORT WORTH|
|  FT WORTH|
+-----+
```

```
df.sort("State").show()
```

```
+-----+-----+-----+-----+
|          City|State|ZipCodeType|Zipcode|
+-----+-----+-----+-----+
|          MESA|  AZ|   STANDARD|  85209|
|          MESA|  AZ|   STANDARD|  85210|
|      HILLIARD|  FL|   STANDARD|  32046|
|    PARC PARQUE| PR|   STANDARD|    704|
|PASEO COSTA DEL SUR| PR|   STANDARD|    704|
|      BDA SAN LUIS| PR|   STANDARD|    709|
|    URB EUGENE RICE| PR|   STANDARD|    704|
| CINGULAR WIRELESS| TX|    UNIQUE|  76166|
|      FORT WORTH| TX|   STANDARD|  76177|
|        FT WORTH| TX|   STANDARD|  76177|
+-----+-----+-----+-----+
```

```
df.sort("City").show()
```


```
+-----+-----+-----+-----+
|          City|State|ZipCodeType|Zipcode|
+-----+-----+-----+-----+
|      BDA SAN LUIS|  PR|   STANDARD|    709|
| CINGULAR WIRELESS| TX|    UNIQUE|  76166|
|      FORT WORTH| TX|   STANDARD|  76177|
|        FT WORTH| TX|   STANDARD|  76177|
|      HILLIARD|  FL|   STANDARD|  32046|
|          MESA|  AZ|   STANDARD|  85209|
|          MESA|  AZ|   STANDARD|  85210|
|    PARC PARQUE|  PR|   STANDARD|    704|
|PASEO COSTA DEL SUR| PR|   STANDARD|    704|
|    URB EUGENE RICE| PR|   STANDARD|    704|
+-----+-----+-----+-----+
```

```
df.describe().show()
```

summary	City	State	ZipCodeType	Zipcode
count	10	10	10	10
mean	null	null	null	43380.6
stddev	null	null	null	39635.30844027274
min	BDA SAN LUIS	AZ	STANDARD	704
max	URB EUGENE RICE	TX	UNIQUE	85210

## 10. PySpark SQL Database

- 1) To Create a Database, first go to the Compute tab and create a new cluster.

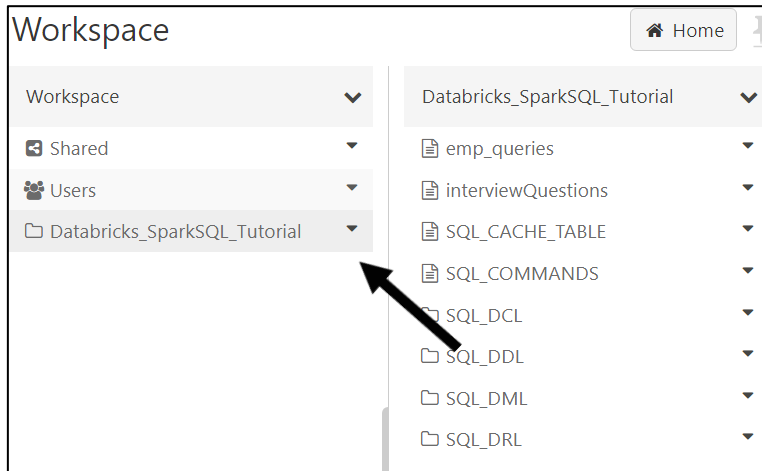


- 2) Give a name for the cluster  
Ex: Data.

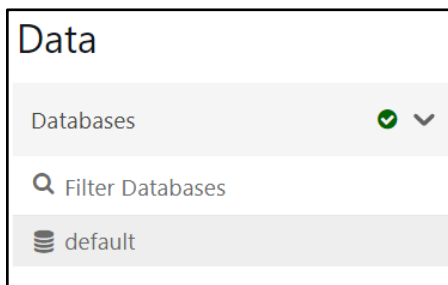
The screenshot shows the 'Create compute' configuration page in Databricks. The breadcrumb 'Compute >' is at the top left. The cluster name 'Data' is entered and has a green checkmark. Below the name are tabs for 'Configuration', 'Notebooks (0)', 'Libraries', 'Event log', and 'Spark UI'. The 'Configuration' tab is active. It contains two main sections: 'Databricks Runtime Version' with a dropdown menu showing '12.2 LTS (includes Apache Spark 3.3.2, Scala 2.12)', and 'Driver type' with a dropdown menu showing 'Community Optimized' and a specification '15.3 GB Memory, 2 Cores, 1 DBU'.

- 3) Download this .dbc file and upload it to the workspace that you are working on.


[https://github.com/raveendratal/PySparkRaveendra/blob/master/Databricks\\_SparkSQL\\_Tutorial.dbc](https://github.com/raveendratal/PySparkRaveendra/blob/master/Databricks_SparkSQL_Tutorial.dbc)



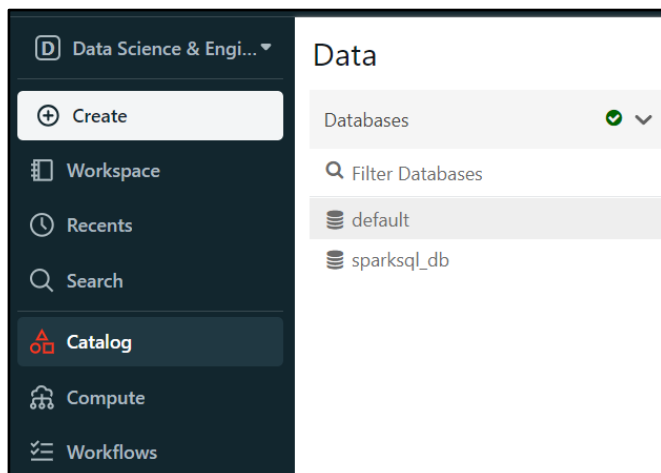
4) You can see the databases inside the cluster.



5) Create a database command.



Then you can see the successfully created database inside the cluster.



```
describe database sparksql_db;
```

Table ▾ +		
	database_description_item ▲	database_description_value ▲
1	Catalog Name	spark_catalog
2	Namespace Name	sparksql_db
3	Comment	
4	Location	dbfs:/user/hive/warehouse/sparksql_db.db
5	Owner	root
↓ 5 rows   0.47 seconds runtime		

```
show databases;
```

1	show databases;
Table ▾ +	
	databaseName ▲
1	default
2	sparksql_db
↓ 2 rows   0.26 seconds runtime	

```
create table sparksql_db.customer(id int, name string)
```

▼ (4) Spark Jobs

- ▶ Job 1 [View](#) (Stages: 1/1)
- ▶ Job 2 [View](#) (Stages: 1/1, 1 skipped)
- ▶ Job 3 [View](#) (Stages: 1/1, 1 skipped)
- ▶ Job 4 [View](#) (Stages: 1/1, 2 skipped)

OK



🔍 Filter Databases

default

sparksql\_db

🔍 Filter Tables

customer ▼

Table Details

sparksql\_db.customer | Refresh

BigData | v

Details

History

Description:

Created at: 2024-03-28 16:34:23

Last modified: 2024-03-28 16:34:25

Partition columns:

Number of files: 0

Size: 0 B

Schema:

	col_name	data_type	comment
1	id	int	null
2	name	string	null

Sample Data:

Table is empty.

```
describe sparksql_db.customer;
```

Table v +

	col_name	data_type	comment
1	id	int	null
2	name	string	null

2 rows

0.26 seconds runtime

```
insert into sparksql_db.customer values(1, 'Ravi')
```

Table v +

	num_affected_rows	num_inserted_rows
1	1	1

1 row

9.09 seconds runtime

Refreshed now



```
select * from sparksql_db.customer;
```

1

select.\* from sparksql\_db.customer1;

▶ (2) Spark Jobs

Table ▼ +

	id	name	
1	1	Ravi	

```
%fs ls /user/hive/warehouse/sparksql_db.db/customer
```

Table ▼ +		
	path	name
1	dbfs:/user/hive/warehouse/sparksql_db.db/customer/_delta_log/	_delta_log/
2	dbfs:/user/hive/warehouse/sparksql_db.db/customer/part-00000-0a89d0eb-d4d3-483a-9d4e-8bbf2d902b9c-c000.snappy.parquet	part-00000-0a89d0eb-d4d3-483a-9d4e-8bbf2d902b9c-c000.sn

```
%python  
spark.conf.get("spark.sql.warehouse.dir")
```

Default warehouse location.

```
1 %python  
2 spark.conf.get("spark.sql.warehouse.dir")  
  
Out[1]: 'dbfs:/user/hive/warehouse'
```

## Conclusion

In conclusion, Spark SQL is a module of Spark that analyses the structured data. It provides Scalability, it ensures high compatibility of the system. Thus, it provides the most natural way to express the Structured Data.